

each cell is scored on two kinds of question

the sweep

18 datasets

5 backbones

7 attributors

2 splits
random · scaffold
(Bemis–Murcko)



one cell

(dataset, backbone,
attributor, split)

186 audited



Does the attribution describe the MODEL?

occlusion faithfulness · coherence · stability

calibration linkage · confidence/correctness regime

Fidelity $\pm$, sparsity, characterisation score

measurable on all 186 cells



Does it describe the CHEMISTRY?

ground-truth localisation against the labelled
substructure (motif-native, RDKit)

measurable on 142 · unmeasurable on 44

The two come apart: ranking attributors by a model-side score picks the wrong one in 26 of 30 tests.