

## THE SWEEP

18 datasets

5 backbones

7 attributors

2 splits

random · scaffold  
(Bemis–Murcko)



### one cell

(dataset, backbone,  
attributor, split)

186 audited



### Does the attribution describe the model?

occlusion faithfulness, motif coherence, checkpoint stability,  
calibration linkage, regime stratification, Fidelity+ / –

measurable on every one of the 186 cells



### Does it describe the chemistry?

ground-truth localisation against the labelled substructure,  
decomposed into RDKit motifs

measurable on 142 cells · no per-atom labels for the other 44

The two answers disagree: ranking attributors by a model-side score picks the wrong one in 26 of 30 tests.